



# Image Classification on Tiny ImageNet



CS 172B NEURAL NETWORKS AND DEEP LEARNING PROJECT

FRANK DONG, NICHOLAS PAVINI, TIANYUN YUAN, JINGQI YAO, NIMA HENDI


# Table of contents

- 1 → INTRODUCTION
- 2 → DATASET
- 3 → MACHINE LEARNING MODELINGS
- 4 → CONCLUSION AND DISCUSSION
- 5 → CONTRIBUTIONS

1

# INTRODUCTION

The ImageNet Large Scale Visual Recognition Challenge(ILSVRC) started in 2010 and has become the standard benchmark of image recognition. Tiny ImageNet Challenge is a similar challenge with a smaller dataset and less image classes.



# Previous Work

- **AlexNet (2012)**
  - 60 MILLION PARAMETERS AND 65,000 NEURONS, 5 CONVOLUTIONAL LAYERS
  - TOP-5 ERROR RATE OF 18.9% IN THE ILSVRC-2010 CONTEST
  - USES DATA AUGMENTATION AND DROPOUT TO FIGHT OVERFITTING
- **VGGNet (2014)**
  - DEEPER CNN WITH SMALLER FILTERS
  - TOP-5 TEST ERROR OF 6.8% IN ILSVRC-2014 CONTEST
  - HAVE DEGRADATION PROBLEM
- **ResNet (2015)**
  - PREVENT DEGRADATION OF DEEPER CNNs



# Project Overview

- Distinguish performance difference between different image classification systems using regular neural networks, classic CNNs, Multi-Layer CNN, VGGNet16 and ResNet
- Apply data augmentation, dropout to prevent overfitting of our models
- Use transfer learning, layer-wise fine-tuning to initialize weight
- Present data visualization of performance between different models


2

# DATASET

# What is Tiny ImageNet?

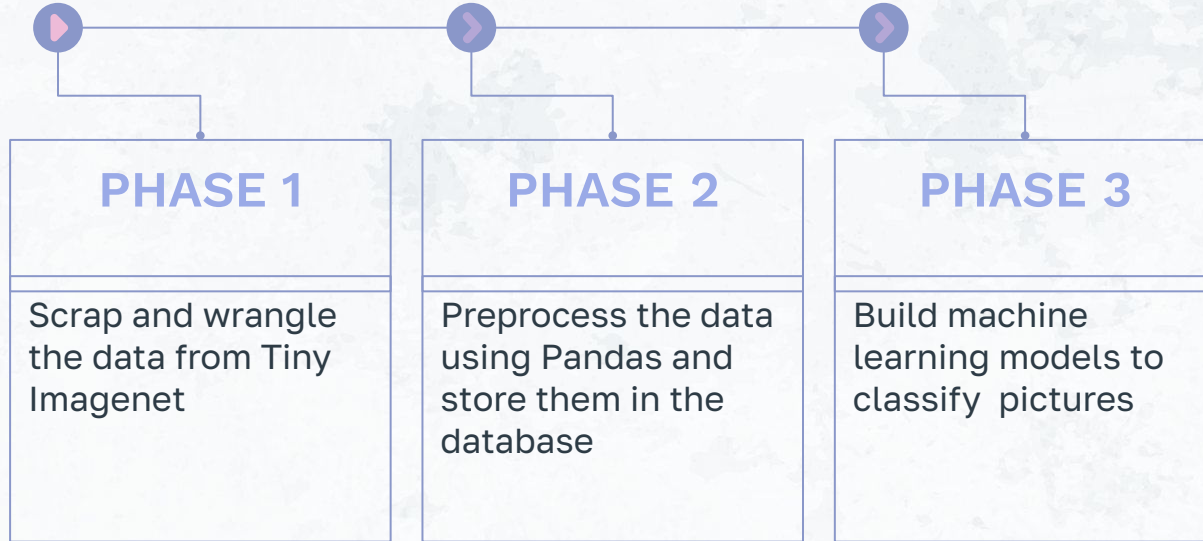
**[define]** Tiny ImageNet is a dataset containing 100,000 images of 200 classes (500 for each class) downsized to  $64 \times 64$  colored images. Each class has 500 training images, 50 validation images and 50 test images. It serves as a benchmark for evaluating image classification models.

**[use]** Bigger and more challenging datasets can be modeled using Tiny ImageNet since it is easier to model.





# Data Processing




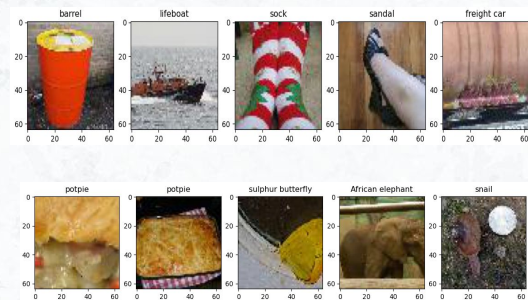
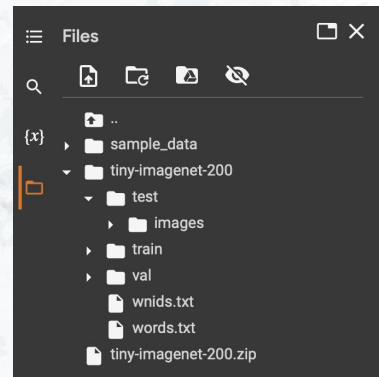
# Tiny ImageNet Dataset

## Data PreProcess

We applied image transformation:

- Convert Images to tensors to process inside the network
- Create the dataset and the DataLoaders to benefit from batching and shuffling
- Change the directory tree to meeting the requirements of the data loader

image (image)	label (class label)
	0 (n01443537)
	0 (n01443537)
	0 (n01443537)
	0 (n01443537)
	0 (n01443537)
	0 (n01443537)
	0 (n01443537)



# Tiny ImageNet Dataset

## Data Augmentation

To prevent overfitting, we augment the data by horizontal flipping, translation and rotation.



original picture



Contrast Adjustment 1.3



Contrast Adjustment 0.7



Horizontal Flipping



Random Rotation



Random Translation

# PROJECT CHALLENGES



**Increased Class Variability:** the Tiny ImageNet dataset contains 200 diverse classes.



**Model Size and Complexity:** The resource constraints of low-power devices necessitate the development of compact models



**Model Efficiency:** achieving high accuracy while considering limited computational resources



3

# MACHINE LEARNING MODELINGS

# LINEAR MODEL


Linear models use only Linear transformations and work best for finding linear correlations in data.

They assume a direct relationship between two or more variables.



Areas of application:

- Finance
- Biology

Typically not useful for image recognition.



# LINEAR MODEL - Architecture and Validation



# LINEAR MODEL - Conclusion

The basic Linear Model had a very low validation accuracy and a much higher training accuracy.

This means that the model was overfitting to the training data and not finding the correlations it needed to work on images outside of the training dataset.

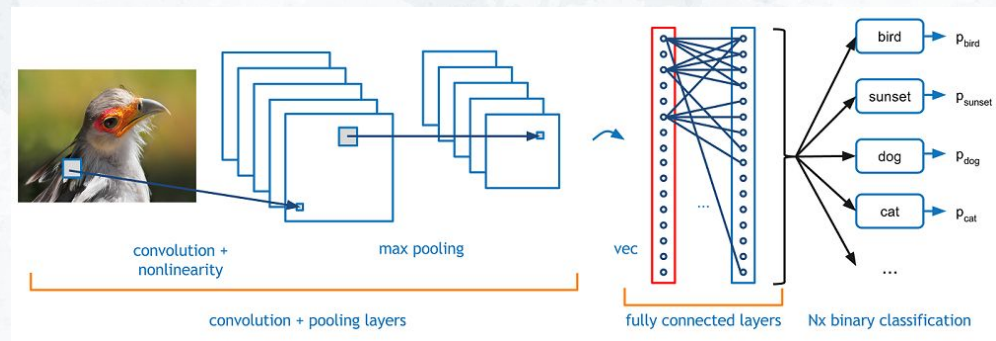
- Final Training Accuracy: 44.9%
- Final Validation Accuracy: 8.38%



# CNN

CNNs are specialized neural network architectures designed to process visual data by mimicking the organization of the human visual cortex.


CNNs use convolutional layers to extract local features, pooling layers to downsample and aggregate information, and fully connected layers for classification.



# 2-LAYER CNN

We used a basic 2-layer CNN (kernel size=3, stride=1, padding=1) with maxpool and dropouts. Because of the convolutional layers' ability to spot features regardless of its position, the model can achieve high training accuracy. However, similar to linear model, basic CNN model also has the problem of overfitting.

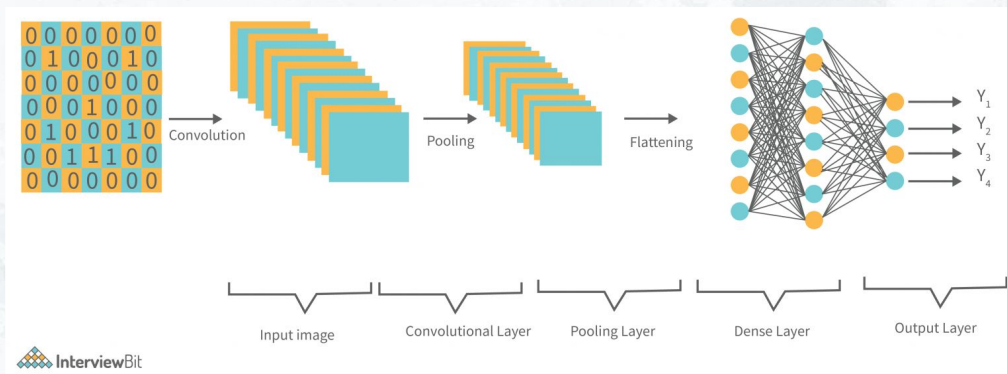
- Final Training Accuracy: 95.2%
- Final Validation Accuracy: 14.9%



# MULTI-LAYER CNN

We used a Multi-layer CNN with maxpool and dropouts

- Final Training Accuracy: 96.1%
- Final Validation Accuracy: 16.3%



InterviewBit

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu',
                           input_shape=(64, 64, 3)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(200, activation='softmax')
])
```

# VGGNet

VGGNet, also known as the Visual Geometry Group Network, is a deep convolutional neural network (CNN) architecture

Simplicity, uniformity, deep structure

3x3 convolutional filters

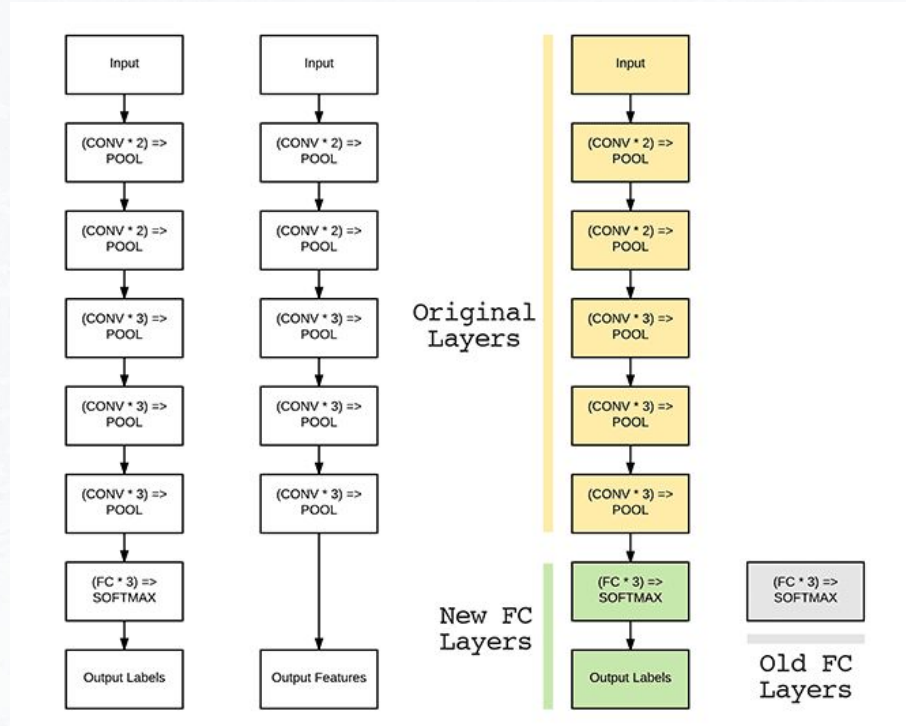
VGGNet has several variations, denoted by the number of layers it contains. VGG16 and VGG19

Model: "vgg16"

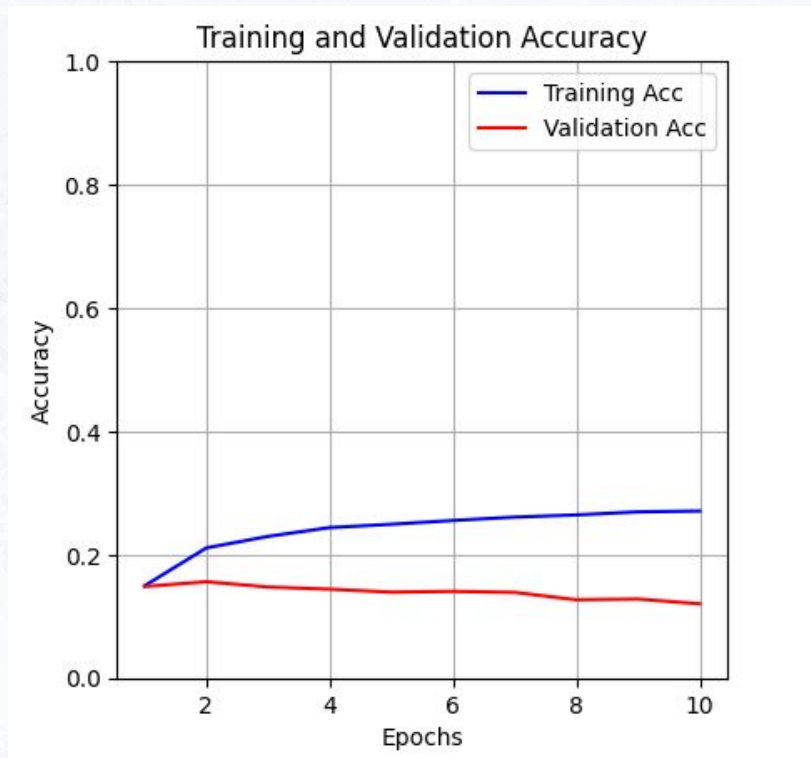
Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 64, 64, 3)]	0
block1_conv1 (Conv2D)	(None, 64, 64, 64)	1792
block1_conv2 (Conv2D)	(None, 64, 64, 64)	36928
block1_pool (MaxPooling2D)	(None, 32, 32, 64)	0
block2_conv1 (Conv2D)	(None, 32, 32, 128)	73856
block2_conv2 (Conv2D)	(None, 32, 32, 128)	147584
block2_pool (MaxPooling2D)	(None, 16, 16, 128)	0
block3_conv1 (Conv2D)	(None, 16, 16, 256)	295168
block3_conv2 (Conv2D)	(None, 16, 16, 256)	590080
block3_conv3 (Conv2D)	(None, 16, 16, 256)	590080
block3_pool (MaxPooling2D)	(None, 8, 8, 256)	0
block4_conv1 (Conv2D)	(None, 8, 8, 512)	1180160
block4_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block4_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block4_pool (MaxPooling2D)	(None, 4, 4, 512)	0
block5_conv1 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block5_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block5_pool (MaxPooling2D)	(None, 2, 2, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		



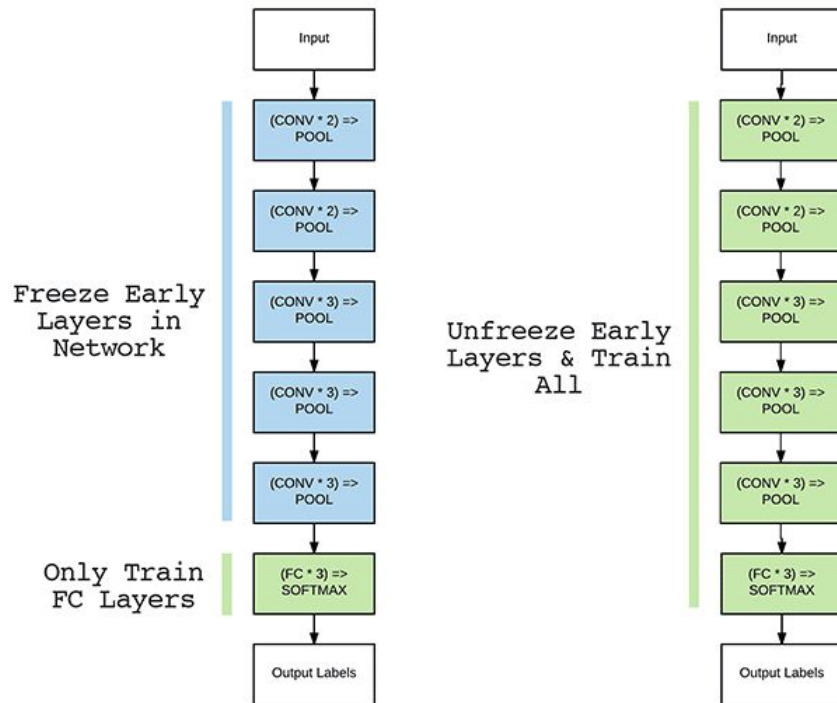
# Transfer Learning





# Transfer Learning



# Layer-wise Fine Tuning



# Layer-wise Fine Tuning


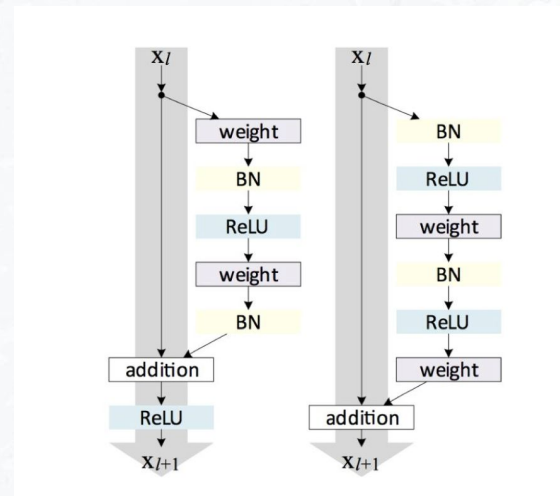




# ResNet


ResNet, short for Residual Neural Network, is a groundbreaking deep learning architecture that revolutionized image classification tasks.

The key innovation lies in the residual blocks, which allow the network to learn residual mappings by directly adding skip connections between layers. These connections mitigate the vanishing gradient problem and facilitate the flow of information, enabling the network to effectively learn complex representations.




# ResNet AND CNN

1. Convolutional layer
2. 2 Residual layer
3. Dense layer

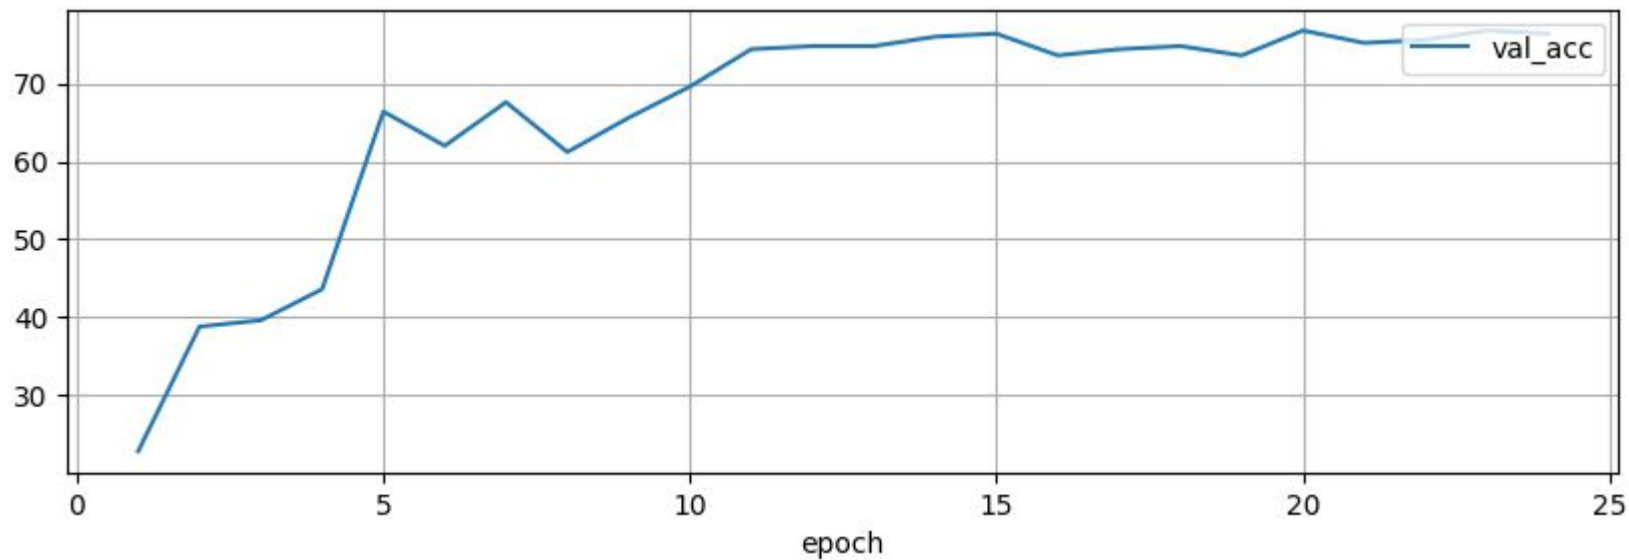


# TRAIN LOSS

[ResNet + CNN]



## [VALIDATION] ResNet + CNN





4

# CONCLUSION AND DISCUSSION



# Why CNN Structure?

- Linear model only looks at pixel level
- Local Connectivity and weight sharing
- RGB image  $\rightarrow$  3 in\_channels



# ResNet Performs Better Than Classic CNN

- More layers (152 vs 2)
- Skip connections
  - addresses the vanishing gradient problem
- Batch normalization
  - Stabilizes and regularizes the training process





Figure 2. Residual learning: a building block.




5

# CONTRIBUTIONS

# OUR TEAM

- (-) Frank Dong  
Responsible for slides structure, ResNet + CNN Model training, validation and report
- (-) Nicholas Pavini  
Responsible for Data Parsing and visualization, and Linear Model training / validation
- (-) TianYun Yuan  
Data Visualization, CNN and regular neural network comparison, VGGNet 16 training/validation, Data preprocessing
- (-) Jingqi Yao  
CNN model training, experimented and compare with ResNet
- (-) Nima Hendi:  
Responsible for Data Wrangling, Multi-Layer CNN Model, Validation and Restructuring of slides and the report



# Reference

A. Krizhevsky, I. Sutskever, and G. E. Hinton. Imagenet classification with deep convolutional neural networks. Pages 1097–1105, 2012.

S. Ioffe and C. Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. arXiv preprint arXiv:1502.03167, 2015.

K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. arXiv preprint arXiv:1409.1556, 2014

N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. The Journal of Machine Learning Research, 15(1):1929–1958, 2014.